

Question 36: Hypothesis Space Search in a Genetic Algorithm Evolving Neural Network Weights

SDG 2 – Zero Hunger

Name: Vishwa V | Reg No: 192421374

Question: Explain the concept of hypothesis space search in the context of a genetic algorithm evolving neural network weights for predicting famine risk in vulnerable regions.

1. What is a Hypothesis Space?

In machine learning, a hypothesis is a candidate model that maps inputs to outputs – here, a neural network with a specific set of weights that maps famine-risk indicators (rainfall, crop yield, food prices, conflict indicators, displacement data, etc.) to a predicted famine-risk score. The hypothesis space is the set of all possible hypotheses the chosen model architecture can represent. For a neural network with a fixed structure (fixed number of layers and neurons), the hypothesis space is the set of all possible weight and bias configurations – effectively an enormous, high-dimensional, continuous space of real-valued vectors.

2. Hypothesis Space Search Using a Genetic Algorithm

Instead of using gradient descent to move through this space one small step at a time, a genetic algorithm treats the search for good weights as an evolutionary search across the hypothesis space:

- **Individual = Hypothesis:** each chromosome encodes one complete weight vector (all weights and biases of the network flattened into a single array) – i.e., one specific point in the hypothesis space.
- **Population = Region of the space explored:** a population of many such weight vectors represents many different points scattered across the hypothesis space at once, rather than a single point as in gradient descent.
- **Fitness = Quality of hypothesis:** each individual's fitness is evaluated by feeding historical regional data through the network with those weights and measuring how accurately it predicts known famine-risk outcomes (e.g., inverse of prediction error, or an F1-score if risk is treated as a classification).
- **Selection, crossover, mutation = Search operators:** these operators move the search through the hypothesis space – selection favors promising regions (weight vectors with high fitness), crossover combines partial solutions from two good hypotheses to explore new combinations, and mutation perturbs weights slightly to explore neighbouring points and escape local optima.

3. Why Frame It as a Search Problem?

Viewing weight optimization as hypothesis space search highlights why a GA is useful here: the famine-risk prediction landscape (mapping weight vectors to prediction accuracy) is likely non-convex, noisy, and may have many local optima due to complex, non-linear relationships

between socio-economic and climatic factors. A GA does not require the landscape to be differentiable and can maintain multiple candidate hypotheses in parallel, making it more robust to getting stuck in poor local optima compared to purely gradient-based search, especially when the fitness signal is derived from limited or noisy real-world humanitarian data.

4. Guiding the Search Toward Good Hypotheses

- **Exploration vs. exploitation:** mutation and random initialization maintain diversity (exploration of the hypothesis space), while selection and elitism concentrate the population around high-fitness hypotheses (exploitation of promising regions).
- **Generational improvement:** across generations, the population's average fitness increases as poor hypotheses (weight vectors that predict famine risk badly) are replaced by offspring of better-performing ones, effectively narrowing the search toward high-accuracy regions of the hypothesis space.
- **Avoiding overfitting:** fitness should be evaluated on a validation subset of regional data separate from training data, so the search converges toward hypotheses that generalize to new vulnerable regions rather than memorizing historical cases.

5. Link to SDG 2 (Zero Hunger)

By searching the hypothesis space for network weights that most accurately predict famine risk, the GA produces an early-warning model that can help humanitarian organizations and governments identify vulnerable regions before a crisis fully develops, enabling earlier intervention – directly supporting SDG 2's goal of ending hunger and ensuring food security, particularly for populations most at risk.